

DialDefer: A Framework for Detecting and Mitigating LLM Dialogic Deference

Parisa Rabbani, Priyam Sahoo, Ruben Mathew, Aishee Mondal,
Harshita Ketharaman, Nimet Beyza Bozdog, Dilek Hakkani-Tür

University of Illinois Urbana-Champaign

{rabbani8, priyams3, rubenom2, aisheem2, hk76, nbozdog2, dilek}@illinois.edu

Abstract

LLMs are increasingly used as third-party judges, yet their reliability when evaluating speakers in dialogue remains poorly understood. We show that LLMs judge identical claims differently depending on framing: the same content elicits different verdicts when presented as a statement to verify (“*Is this statement correct?*”) versus attributed to a speaker (“*Is this speaker correct?*”). We call this *dialogic deference* and introduce DIALDEFER, a framework for detecting and mitigating these framing-induced judgment shifts. Our Dialogic Deference Score (DDS) captures directional shifts that aggregate accuracy obscures. Across nine domains, 3k+ instances, and four models, conversational framing induces large shifts (IDDSI up to 87pp, $p < .0001$) while accuracy remains stable (< 2 pp), with effects amplifying 2–4 $\times$ on naturalistic Reddit conversations. Models can shift toward agreement (deference) or disagreement (skepticism) depending on domain — the same model ranges from DDS = -53 on graduate-level science to $+58$ on social judgment. Ablations reveal that human-vs-LLM attribution drives the largest shifts (17.7pp swing), suggesting models treat disagreement with humans as more costly than with AI. Mitigation attempts reduce deference but can over-correct into skepticism, framing this as a calibration problem beyond accuracy optimization.¹

1 Introduction

LLMs are increasingly deployed as third-party judges: evaluating response quality (Zheng et al., 2023; Liu et al., 2023), moderating content (Kolla et al., 2024), and arbitrating interpersonal disputes (Susskind, 2023). Online forums like Reddit’s r/AmIOverreacting² host millions of users posting private conversations to ask strangers

¹Code and Data available at <https://github.com/LadyPary/DialDefer>

²<https://www.reddit.com/r/AmIOverreacting/>




Figure 1: **Task framing flips judgment on identical content.** Under **factual inquiry** (C_1 ; left), GPT-4o-mini correctly rejects the claim. Under **conversational judgment** (C_2 ; right), the model endorses the same claim attributed to a speaker. This directional shift toward agreement is *dialogic deference*.

whether someone is overreacting. Given this demand, many users now turn to LLMs for similar social arbitration. But this use case is high-risk. Alignment methods designed for helpfulness can conflict with calibrated judgment (Sharma et al., 2024; Perez et al., 2023), and real-world incidents already document models validating delusional beliefs (Editorial, 2025; Preda, 2025) and reinforcing harmful ideation (Schoene and Canca, 2025; Rust and Chang, 2025). Prior work on sycophancy documents how models over-accommodate user viewpoints at the cost of accuracy (Sharma et al., 2024; Cheng et al., 2025b; Hong et al., 2025). However, these studies cast the model as a direct conversational partner responding to an identifiable user. Much less is known about what happens when the model serves as a third-party judge evaluating others’ exchanges where no explicit “user” exists to

please. Initial evidence suggests dialogue framing can flip judgments (Rabbani et al., 2026), but scope, mechanisms, and mitigations remain unexplored.

This paper introduces DIALDEFER, a framework for detecting and mitigating *Dialogic Deference*. Following Rabbani et al. (2026), we contrast two conditions: factual inquiry (C_1 : “Is this statement correct?”) versus conversational judgment (C_2 : “Is this speaker correct?”), holding semantic content fixed while varying only framing. As Figure 1 illustrates, a model may correctly reject a claim under C_1 but endorse the identical claim when attributed to a speaker in C_2 .

A key challenge is that standard accuracy can obscure dialogic deference entirely. A model that becomes more agreeable under conversational framing endorses both correct and incorrect speakers more readily; accuracy increases on the former but decreases on the latter. If these shifts balance, average accuracy remains stable while behavior shifts dramatically. To capture this, we introduce the **Dialogic Deference Score (DDS)**, measuring whether models become disproportionately deferential ($\text{DDS} > 0$) or skeptical ($\text{DDS} < 0$). The ideal is high accuracy with $\text{DDS} \approx 0$. We observe that DDS values range from -53pp to $+87\text{pp}$ across models and domains.

Across a Unified Benchmark spanning nine domains, 3,244 items, and four models (GPT-4o, GPT-4o-mini, Gemma-3-12B, Qwen-2.5-7B), conversational framing produces large directional shifts: models maintain stable average accuracy ($< 2\text{pp}$ change) while exhibiting DDS exceeding $+30$. Effects are domain-dependent: GPT-4o ranges from $\text{DDS} = -53$ on GPQA (graduate-level science) to $+58$ on r/AIO (social judgment). On r/AIO ($N=280$), a naturalistic dataset of real multi-turn interpersonal conflicts, every model shows its highest DDS, with Gemma-3-12B reaching $+87$ versus $+29$ on synthetic benchmarks, confirming that laboratory findings underestimate real-world susceptibility.

To understand why dialogic deference occurs, we conduct speaker-label ablations and analyze 2,414 judgment flips. Human-vs-LLM attribution produces the largest shifts (17.7pp swing), while demographic cues have minimal effect, suggesting models treat disagreement with humans as costlier than with AI. Reasoning analysis identifies two dominant failure modes: internal incoherence (30%), where C_2 reasoning contradicts C_1 conclusions, and social framing (27%), where models

validate feelings rather than evaluate facts. Prompt-based interventions reduce DDS by $10\text{--}23\text{pp}$ without accuracy gains; SFT improves C_2 accuracy by 22pp and reduces DDS by 24pp . However, interventions often over-correct into skepticism, framing dialogic deference as fundamentally a calibration problem.

Contributions. (1) DIALDEFER, a framework for isolating dialogic framing effects via paired conditions, with a Unified Benchmark spanning nine domains validated on naturalistic r/AIO data (§3). (2) The Dialogic Deference Score (DDS), revealing systematic judgment shifts that aggregate accuracy obscures (§4). (3) Ablation and mitigation experiments showing human-vs-AI attribution drives deference, smaller models are most susceptible, and reducing deference without over-correction remains open (§5).

2 Related Work

LLM Reliability Issues. LLMs exhibit reliability failures including strategic deception under pressure (Scheurer et al., 2023; Meinke et al., 2024) and sycophancy, where models prioritize user agreement over accuracy (Perez et al., 2023; Sharma et al., 2024) that are behaviors arising from alignment techniques like RLHF (Ouyang et al., 2022; Wei et al., 2024). The MASK benchmark confirms that accuracy and honesty are dissociable (Ren et al., 2025). Research has evolved from single-turn sycophancy (Perez et al., 2023) to multi-turn robustness benchmarks (Hong et al., 2025; Fanous et al., 2025) and social sycophancy in narrated scenarios (Cheng et al., 2025b). However, prior work assumes the model is a direct conversational participant, and cannot explain judgment shifts in third-party evaluation where no explicit “user” exists to please. Moreover, existing metrics focus on accuracy degradation, missing the complementary phenomenon DDS captures: apparent accuracy gains from increased agreeableness. Preliminary evidence suggests dialogue framing can flip judgments (Rabbani et al., 2026), but scope and mechanisms remain unexplored.

LLM as a Third-Party Judge. LLM-based evaluation offers scalable alternatives to surface-level metrics, with G-Eval (Liu et al., 2023) and MT-Bench (Zheng et al., 2023) achieving strong human correlation. Known biases include length, verbosity, and position (Dubois et al., 2024; Liu




Figure 2: **The DialDefer framework** transforms datasets into paired experimental conditions. **(a) Unified Benchmark:** Question–answer pairs become Factual Inquiry prompts (C_1 : “Is this statement correct?”) vs. Conversational Judgment prompts (C_2 : “Is Speaker 2 correct?”). **(b) r/AIO:** Reddit conversations are neutralized (first-person pronouns replaced with speaker labels). Speaker 1 is the other party in the conversation; Speaker 2 is the original poster (OP). We construct a judgment about Speaker 1 (“Speaker 1 is [NOT] overreacting”) and test whether attributing it to Speaker 2 changes model evaluation. Ground truth is derived from community consensus (inverted, since the community judges the OP). **Right:** Paired conditions share identical content: $C_1^T \leftrightarrow C_2^C$ (correct content) and $C_1^F \leftrightarrow C_2^I$ (incorrect content). The Dialogic Deference Score captures framing-induced judgment shift: $\text{DDS} = \Delta_{\text{Correct}} - \Delta_{\text{Incorrect}}$, where $\text{DDS} > 0$ indicates deference and $\text{DDS} < 0$ indicates skepticism.

et al., 2024), but a more fundamental vulnerability remains: whether social dynamics of evaluated content can bias the judge itself. This parallels psychology’s authority bias, where claims are more persuasive when speaker-asserted (Milgram, 1963; Cialdini, 2001), and speech act theory’s distinction that *who* says something carries weight beyond *what* is said (Searle, 1969). Our findings suggest LLMs exhibit analogous behavior, with implications for content moderation and dispute resolution.

3 Methodology

We introduce **DialDefer**, a framework for measuring how dialogic framing affects LLM judgment. When identical information is presented as a statement to verify versus a speaker to evaluate, models exhibit systematic judgment shifts. We quantify these shifts using the *Dialogic Deference Score* (DDS), a directional metric capturing whether conversational framing induces deference or skepticism.

Figure 2 illustrates the framework. We evalu-

ate two settings: (1) a **Unified Benchmark** comprising nine Q/A datasets transformed into paired conditions (Figure 2a), and (2) **r/AIO**, a naturalistic dataset of interpersonal conflicts from Reddit (Figure 2b), testing whether effects generalize to complex, emotionally charged conversations.

3.1 Experimental Design

Following Rabbani et al. (2026), we construct paired conditions (Figure 2). C_1 : **Factual Inquiry** (“Is this statement correct?”) probes knowledge without speaker attribution; C_2 : **Conversational Judgment** (“Is Speaker X correct?”) introduces dialogic framing while preserving semantic content. Crossing framing with answer correctness yields four cases: C_1 -True/ C_1 -False and C_2 -Correct/ C_2 -Incorrect. Paired conditions share identical content ($C_1^T \leftrightarrow C_2^C$; $C_1^F \leftrightarrow C_2^I$), so any difference in accuracy isolates the effect of dialogic framing.