

that one party being NOR implies the other is OR. Our analysis confirms this holds in the majority of cases, but the assumption may not universally apply; some conflicts involve mutual overreaction or neither party overreacting. Third, the neutralization step in our preprocessing pipeline introduces potential artifacts, though we mitigate this by applying identical processing to both experimental conditions; all dialogue transcripts and best-comment verdict labels were human-verified. Finally, community consensus serves as an imperfect proxy for ground truth in subjective social judgments; reasonable people may disagree about whether someone is “overreacting.” Future work should expand the OR subset and explore alternative ground-truth definitions, such as expert annotations or multi-annotator agreement with inter-rater reliability metrics.

Calibration Tradeoffs in Mitigation. Our mitigation results reveal a fundamental calibration challenge: reducing deference can over-correct into dialogic skepticism. We trained SFT and DPO on benchmark data excluding TruthfulQA and r/AIO to test generalization. On held-out benchmarks, SFT achieves the largest accuracy improvement (+22.0pp) and DDS reduction ($\downarrow$ 24.1), generalizing well to TruthfulQA despite never seeing it during training. However, both fine-tuning methods fail to generalize to r/AIO, exhibiting atypical universal-agreement behavior that increases rather than decreases deference asymmetry. This suggests that naturalistic social judgment requires dedicated training data beyond synthetic benchmarks.

Importantly, our mitigation evaluation measures improvement in the *dialogic* condition only: we compute DDS using a fixed C1 baseline (no mitigation applied) and measure how much C2 accuracy improves under intervention. This design choice means our mitigations target framing sensitivity specifically; they may simultaneously affect C1 performance in ways we do not optimize for. A fully robust solution would require joint calibration across both conditions.

Despite these limitations, fine-tuning consistently outperforms prompting interventions in both accuracy gains and DDS reduction, representing a promising direction for future work. SFT over-corrects on three domains (PlausibleQA, HARP, HaluEval), flipping from positive to negative DDS, and the “Be Honest” prompt shows similar patterns. No single intervention eliminates deference without domain-specific side effects. Future work

should focus on precision calibration: training objectives and data selection strategies that target domain-robust DDS near zero while preserving correctness across *both* factual and conversational conditions, rather than shifting global agree/disagree thresholds.

Model Coverage. We evaluate four models from three families (GPT-4o, GPT-4o-mini, Qwen-2.5-7B-Instruct, Gemma-3-12B-IT), providing coverage across closed-source and open-weight systems, but this represents a small fraction of deployed LLMs. Our finding that smaller open-weight models exhibit stronger deference (Qwen: DDS=+33.8, Gemma: +29.5) than larger closed-source models (GPT-4o: -1.1) warrants replication across additional model families and scales. Our tested models span instruction-tuned (Gemma, Qwen) and RLHF-trained variants (GPT-4o family). The pattern (stronger deference in instruction-tuned models) suggests alignment training aimed at helpfulness may inadvertently strengthen social accommodation, but we cannot disentangle these contributions without controlled comparisons across training paradigms.

8 Ethical Considerations

This work reveals that LLMs exhibit systematic judgment shifts when identical content is framed as dialogue rather than factual inquiry, a vulnerability with implications for deployed systems.

Dual-Use Considerations. Our findings demonstrate that conversational framing alone can shift LLM judgments by over 30 percentage points, with effects amplifying 2–4 $\times$ in naturalistic settings. This vulnerability could theoretically be exploited by adversaries seeking to manipulate LLM-as-judge systems, for instance by strategically framing claims as speaker utterances to increase acceptance rates in content moderation pipelines. However, we believe disclosure serves the research community: without awareness of dialogic deference, developers cannot design robust safeguards. The phenomenon likely emerges from standard alignment training (models treating disagreement with humans as costlier), meaning the vulnerability already exists in deployed systems. Our contribution is diagnostic, not enabling. We note that related effects are already documented in the alignment literature under different framings (sycophancy, social desirability bias); our work systematizes detection

rather than introducing novel attack vectors.

Deployment Risks. Our findings reveal a troubling inverse relationship between deployment prevalence and judgment reliability. The models most attractive for deployment (cheap to run, easy to scale, increasingly used in LLM-as-judge pipelines) exhibit the strongest deference biases. Qwen-2.5-7B (DDS=+33.8) and Gemma-3-12B (+29.5) show deference effects an order of magnitude larger than GPT-4o (−1.1). Furthermore, amplification from synthetic benchmarks to naturalistic settings suggests that laboratory evaluations provide false assurance; systems deployed in real-world conversational contexts face amplified risk that standard evaluations fail to capture.

Data and Privacy. The r/AIO dataset derives from publicly posted Reddit content that users voluntarily shared in a public forum seeking community judgment. We follow Reddit’s API Terms of Service and research guidelines. As detailed in Appendix B, we publicly release only post IDs and fully anonymized processed data; raw Reddit data is available to researchers upon request for reproducibility purposes. No models were trained on this data. The benchmark datasets (TruthfulQA, BBQ, SocialQA, etc.) are established public resources used in accordance with their licenses. Given that the content was already publicly posted by users seeking external opinions, and that our work provides diagnostic tools rather than enabling harm, we believe the research benefit outweighs minimal additional privacy risk.

Mitigation and Responsible Disclosure. We explore multiple mitigation strategies (prompting interventions, SFT, DPO) and report both successes and failures transparently, including cases where interventions over-correct. Our goal is to provide actionable guidance for practitioners, not a definitive solution. We acknowledge that mitigating dialogic deference without inducing over-criticality remains an open challenge. We hope this work contributes to evaluation practices and alignment strategies that prioritize judgment consistency alongside helpfulness.

References

Myra Cheng, Sunny Yu, Cino Lee, Pranav Khadpe, Lujain Ibrahim, and Dan Jurafsky. 2025a. [Elephant: Measuring and understanding social sycophancy in llms](#). *Preprint*, arXiv:2505.13995.

Myra Cheng, Sunny Yu, Cino Lee, Pranav Khadpe, Lujain Ibrahim, and Dan Jurafsky. 2025b. [Social sycophancy: A broader understanding of LLM sycophancy](#). *Preprint*, arXiv:2505.13995v1.

Robert B. Cialdini. 2001. *Influence: Science and Practice*, 4th edition. Allyn & Bacon.

Yann Dubois, Percy Liang, and Tatsunori B. Hashimoto. 2024. Length-controlled AlpacaEval: A simple way to debias automatic evaluators. In *Proceedings of the Conference on Language Modeling*.

Editorial. 2025. [Can AI chatbots validate delusional thinking?](#) *BMJ*.

Aaron Fanous, Jacob Goldberg, Ank A. Agarwal, Joanna Lin, Anson Zhou, Sonnet Xu, Vasiliki Bikia, Roxana Daneshjou, and Sanmi Koyejo. 2025. [SycEval: Evaluating LLM sycophancy](#). In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*, pages 893–900.

Gemma Team. 2025. [Gemma 3 technical report](#). *Preprint*, arXiv:2503.19786.

Olga Golovneva, Moya Chen, Spencer Poff, Martin Corredor, Luke Zettlemoyer, Maryam Fazel-Zarandi, and Asli Celikyilmaz. 2023. [ROSCOE: A suite of metrics for scoring step-by-step reasoning](#). In *Proceedings of the International Conference on Learning Representations*.

Jiseung Hong, Grace Byun, Seungone Kim, Kai Shu, and Jinho D. Choi. 2025. [Measuring sycophancy of language models in multi-turn dialogues](#). In *Findings of the Association for Computational Linguistics: EMNLP 2025*.

Minbeom Kim, Hwanhee Lee, Joonsuk Park, Hwaran Lee, and Kyomin Jung. 2025. [AdvisorQA: Towards helpful and harmless advice-seeking question answering with collective intelligence](#). In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 6545–6565, Albuquerque, New Mexico. Association for Computational Linguistics.

Mahi Kolla, Siddharth Salunkhe, Eshwar Chandrasekharan, and Koustuv Saha. 2024. [LLM-mod: Can large language models assist content moderation?](#) In *Extended Abstracts of the 2024 CHI Conference on Human Factors in Computing Systems*, pages 1–8, Honolulu, HI, USA. Association for Computing Machinery.

Satyapriya Krishna, Andy Zou, Rahul Gupta, Eliot Krzysztof Jones, Nick Winter, Dan Hendrycks, J. Zico Kolter, Matt Fredrikson, and Spyros Matsoukas. 2025. [D-rex: A Benchmark for detecting deceptive reasoning in large language models](#). *Preprint*, arXiv:2509.17938.

- Junyi Li, Xiaoxue Cheng, Wayne Xin Zhao, Jian-Yun Nie, and Ji-Rong Wen. 2023. [HaluEval: A large-scale hallucination evaluation benchmark for large language models](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*.
- Stephanie Lin, Jacob Hilton, and Owain Evans. 2022. [TruthfulQA: Measuring how models mimic human falsehoods](#). In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 3214–3252, Dublin, Ireland. Association for Computational Linguistics.
- Yang Liu, Dan Iter, Yichong Xu, Shuohang Wang, Ruochen Xu, and Chenguang Zhu. 2023. [G-Eval: NLG evaluation using GPT-4 with better human alignment](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 2511–2522, Singapore. Association for Computational Linguistics.
- Yuxuan Liu, Tianchi Yang, Shaohan Huang, Zihan Zhang, Haizhen Huang, Furu Wei, Weiwei Deng, Feng Sun, and Qi Zhang. 2024. [Calibrating LLM-based evaluator](#). In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, pages 2638–2656, Torino, Italia. ELRA and ICCL.
- Alexander Meinke, Bronson Schoen, J  r  my Scheurer, Mikita Balesni, Rusheb Shah, and Marius Hobbhahn. 2024. [Frontier models are capable of in-context scheming](#). *Preprint*, arXiv:2412.04984.
- Stanley Milgram. 1963. Behavioral study of obedience. *The Journal of Abnormal and Social Psychology*, 67(4):371–378.
- Jamshid Mozafari, Abdelrahman Abdallah, Bhawna Piryani, and Adam Jatowt. 2025. [Wrong answers can also be useful: PlausibleQA – a large-scale QA dataset with answer plausibility scores](#). In *Proceedings of the 48th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 3832–3842, New York, NY, USA. Association for Computing Machinery.
- OpenAI. 2024. [GPT-4o system card](#). *Preprint*, arXiv:2410.21276.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul F Christiano, Jan Leike, and Ryan Lowe. 2022. [Training language models to follow instructions with human feedback](#). In *Advances in Neural Information Processing Systems*, volume 35, pages 27730–27744. Curran Associates, Inc.
- Arjun Panickssery, Samuel R. Bowman, and Shi Feng. 2024. [LLM evaluators recognize and favor their own generations](#). In *Advances in Neural Information Processing Systems*, volume 37, pages 68772–68802, Vancouver, Canada.
- Alicia Parrish, Angelica Chen, Nikita Nangia, Vishakh Padmakumar, Jason Phang, Jana Thompson, Phu Mon Htut, and Samuel Bowman. 2022. [BBQ: A hand-built bias benchmark for question answering](#). In *Findings of the Association for Computational Linguistics: ACL 2022*, pages 2086–2105, Dublin, Ireland. Association for Computational Linguistics.
- Ethan Perez, Sam Ringer, Kamile Lukosiute, Karina Nguyen, Edwin Chen, Scott Heiner, Craig Pettit, Catherine Olsson, Sandipan Kundu, Saurav Kadavath, Andy Jones, Anna Chen, Benjamin Mann, Brian Israel, Bryan Seethor, Cameron McKinnon, Christopher Olah, Da Yan, Daniela Amodei, and 44 others. 2023. [Discovering language model behaviors with model-written evaluations](#). In *Findings of the Association for Computational Linguistics: ACL 2023*, pages 13387–13434, Toronto, Canada. Association for Computational Linguistics.
- Adrian Preda. 2025. [AI-induced psychosis: A new frontier in mental health](#). *Psychiatric News*.
- Parisa Rabbani, Nimet Beyza Bozdog, and Dilek Hakkani-T  r. 2026. From fact to judgment: Investigating the impact of task framing on LLM conviction in dialogue systems. In *Proceedings of the International Workshop on Spoken Dialogue Systems (IWSDS)*, to appear.
- David Rein, Betty Li Hou, Asa Cooper Stickland, Jackson Petty, Richard Yuanzhe Pang, Julien Dirani, Julian Michael, and Samuel R. Bowman. 2024. [GPQA: A graduate-level Google-proof Q&A benchmark](#). In *Proceedings of the First Conference on Language Modeling*.
- Richard Ren, Arunim Agarwal, Mantas Mazeika, Cristina Menghini, Robert Vacareanu, Brad Kenstler, Mick Yang, Isabelle Barrass, Alice Gatti, Xuwang Yin, Eduardo Trevino, Matias G  ralnik, Adam Khoja, Dean Lee, Summer Yue, and Dan Hendrycks. 2025. [The MASK benchmark: Disentangling honesty from accuracy in AI systems](#). *Preprint*, arXiv:2503.03750.
- Susanne Rust and Justin Chang. 2025. [AIs gave scarily specific self-harm advice to users expressing suicidal intent, researchers find](#). *Los Angeles Times*.
- Maarten Sap, Hannah Rashkin, Derek Chen, Ronan Le Bras, and Yejin Choi. 2019. [Social IQa: Commonsense reasoning about social interactions](#). In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 4463–4473, Hong Kong, China. Association for Computational Linguistics.
- J  r  my Scheurer, Mikita Balesni, and Marius Hobbhahn. 2023. [Large language models can strategically deceive their users when put under pressure](#). *Preprint*, arXiv:2311.07590.